
Text-generation inference

Nicolas Patry - 20/03/2023

Overview

- **Optimizing Bloom**
 - TP vs PP
 - Transformers
- **Optimizing the webserver**
 - Latency vs Throughput
 - Minimizing latency AND maximizing throughput.
 - Generation with PastKeyValues
 - Testing

What is bloom ?

- Very large language model (LLM), of 170 billion parameters, on par with GPT-3
 - Trained on many languages (13)
 - Open sourced (including training, datasets and many other artifacts)
 - Requires 340 GB of VRAM (8A100 80, or 16A100 40)
-

An inference server ?

- Enable users to try out bloom without any hardware with an API
 - Optimize for user perceived latency
 - Give the best possible showcase for the model
-

Optimizing Bloom

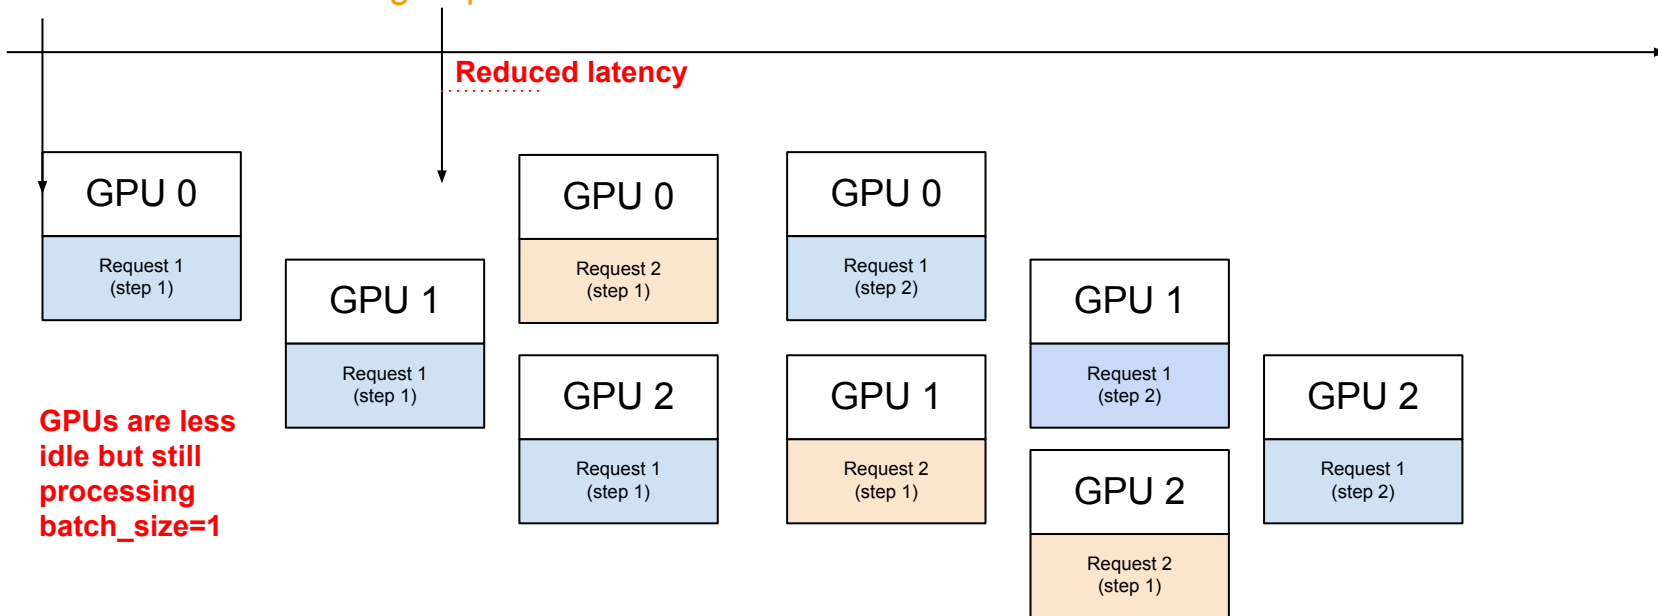
Bloom Optimization

- Start from a CORRECT program version
 - Think early about interactions between core optimizations:
 - Past key values
 - TP vs PP
 - Webserver orchestration
 - Think about your endgoal first. Latency vs throughput, and targets
 - Back of the envelope calculations are critical. How far are you from optimal code on given hardware.
 - Can it be done ? And how hard will it be ?
 - Measure performance
 - Identify bottlenecks
 - Make bottlenecks faster
 - Measure performance, and check it worked
-

TP vs PP

Incoming request 1

Incoming request 2



Past key values

I		I	o	v	e		p	u	p	p	i	e	s
---	--	---	---	---	---	--	---	---	---	---	---	---	---

Tokenization

12	1250	3214	2931
----	------	------	------

Embeddings

1xNx768 Matrix



Attention (simplified)

768x768



1xNx768 Matrix

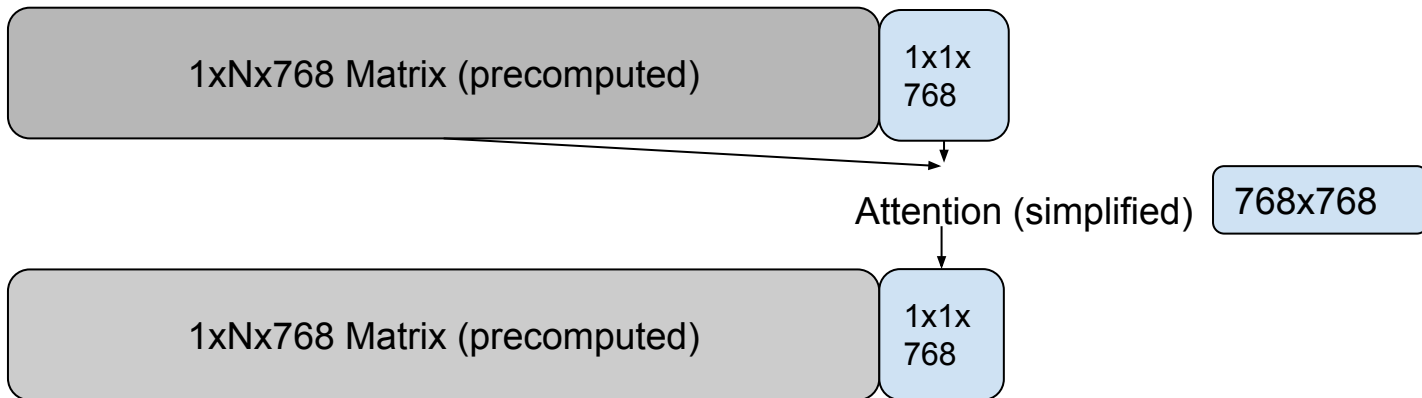
Past key values

I		I	o	v	e		p	u	p	p	i	e	s		!
---	--	---	---	---	---	--	---	---	---	---	---	---	---	--	---


Tokenization

12	1250	3214	2931	367
----	------	------	------	-----

Embeddings



TP vs PP



(a) MLP

TP vs PP

PP

- Simple modeling
- Little GPU communication
- Good for throughput

TP

- More complex modeling (transformers is quite easy)
 - Lots of inter GPU communication (new bottleneck)
 - Good for latency
-

Transformers optimization

- Start by profiling:
https://github.com/pytorch/kineto/blob/main/tb_plugin/README.md
 - Usually start by removing "stupid ops":
 - Remove unnecessary reshapes, tensor creations
 - Fuse multiple small operations (@torch.jit.script)
 - Aim for maximum GPU utilization, maximum tensor cores utilization
 - Find the new bottleneck (usually attention or MLP in transformers)
 - Create new kernels
-

Transformers optimization

- PP to TP:
Latency 350ms/token -> 100ms/tokens (16A100 40Go)
 - Remove extra reshapes (mostly past key values reshapes):
Latency 100 ms -> 90ms
 - Fuse GELU op (@torch.jit.script)
Latency 90ms -> 80ms
 - Rewrite custom kernel for attention (prevent softmax f32 tensor creation)
Latency 80ms -> 70ms/token
 - Better hardware (8A100 80):
Latency 70ms -> 45ms
-

Transformers optimization

- PP to TP:
Half a day (expert)
 - Remove extra reshapes (mostly past key values reshapes):
2 days
 - Fuse GELU op (@torch.jit.script)
10s
 - Rewrite custom kernel for attention (prevent softmax f32 tensor creation)
1 day (expert)
 - Better hardware (8A100 80):
Half a day (modifying hard coded code around, and checking infra)
-

Big Takeaways

- Optimize of your feedback latency (time it takes to test an idea).
 - Be in control of your tools - Drop Torch or Python if they cause issues (custom kernel, or GIL)
 - Measure, measure, measure
 - Profiling is a way to get insights for slow parts/bottlenecks, NEVER measure anything in profiling mode
-

Optimizing the webserver

Webserver

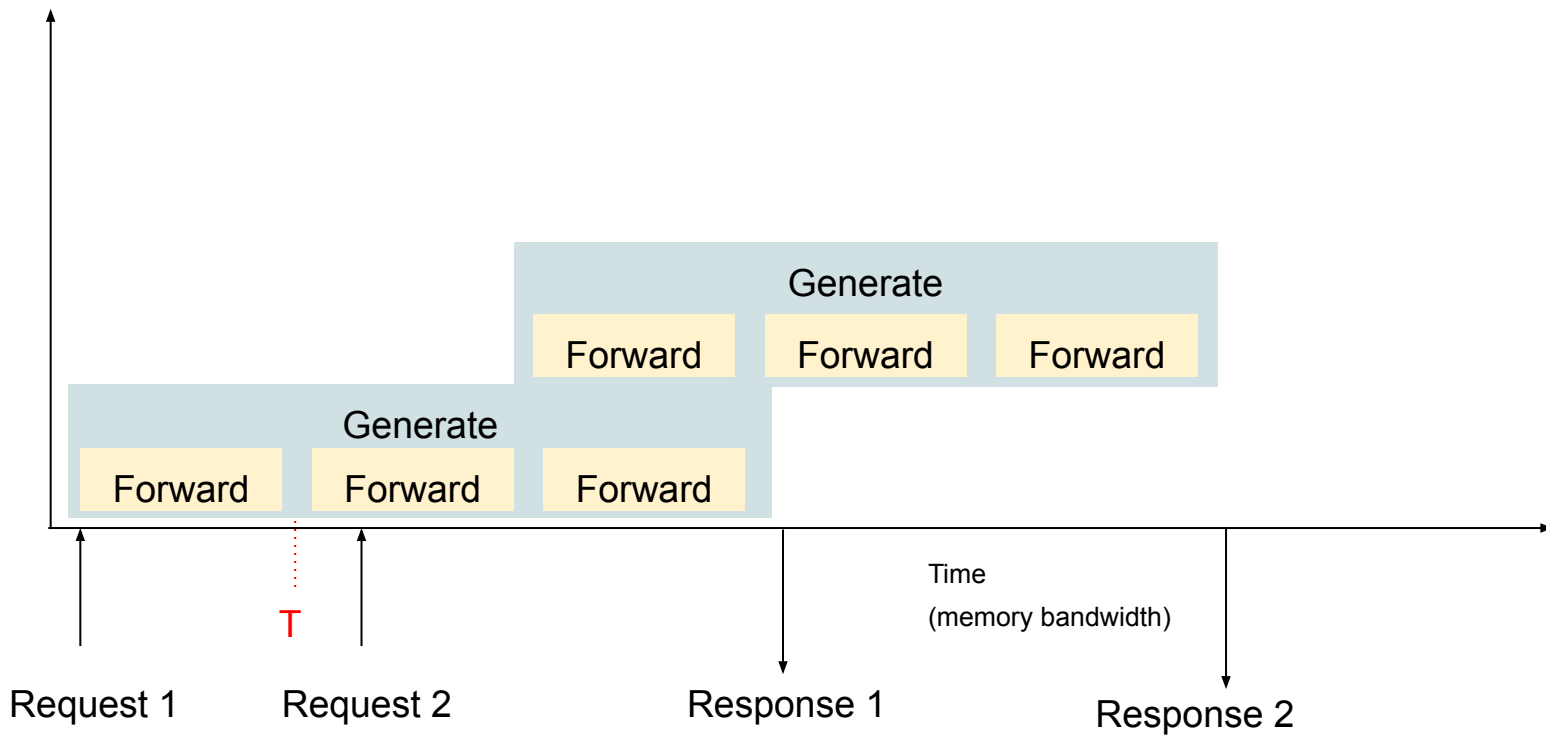
- Lots of incoming requests, at random times
 - Incoming requests have different profiles
 - Different parameters (sampling, greedy, temperature etc..)
 - Different lengths + max_new_tokens
 - We want to minimize latency for EVERYONE
-

Webserver

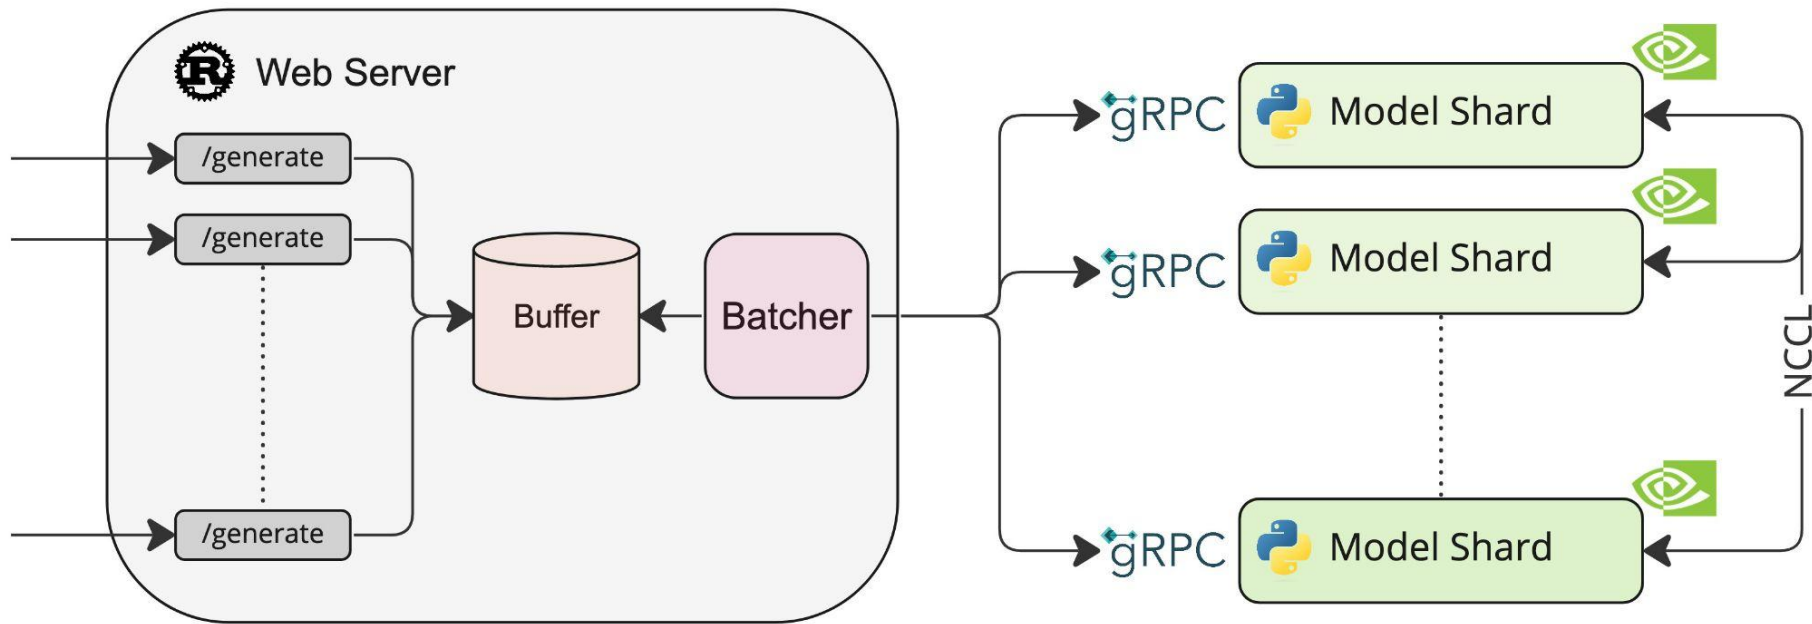
- User perceived performance:
 - 100ms: Instant
 - 1s: Fast
 - 10s: slow
 - 1mn: Too long
-

Webserver

Compute



Webserver



Webserver

- Post process (LogitsProcessor) is much faster than the forward loop, so we can do per request treatment
 - Long requests (`max_new_tokens > 100`) are slow anyway ($>4.5s$) so we can delay it more by allowing small requests (`max_new_tokens < 20`) to enter the batch midway. This requires running a forward WITHOUT past key cache.
 - Every $\sim 1s$ (20tokens) allow small requests to enter, delays large requests by $\sim 5\%$ (more in practice) but it CAPS the overall latency of small requests.
 - Streaming new tokens allows for even better UX, time to something happening is reduced to 500ms (average delay)
-

Big Takeaways

- Load times were the developer bottleneck (10mn to load the weights)
 - Access to big machines is sparse (try code correctness on smaller models *before* testing on larger models)
 - Be in control of your tools - Drop Torch or Python if they cause issues (custom kernel, or GIL)
 - Measure, measure, measure
 - Have clear targets in minds, keep thinking about what you are optimizing for.
-